

Does the Voice Change the Grade?

A Causal Audit of Speaker-Identity Bias in Audio-LLM Judges

Anonymous authors
Paper under double-blind review

Abstract

Audio language models now grade spoken answers directly: automated oral exams, AI interview screening, and speech leaderboards all use an audio-input model as the judge. Every one of these systems assumes the judge scores what was said, not how it sounded. We test that assumption. We hold a bank of 24 interview-style answers fixed and render each one in four accents (American, British, Indian, and Nigerian English) and two speaker genders using text-to-speech, then score every clip with two audio-input judges, Gemini 3.1 Pro and Gemini 3.6 Flash, one clip per call, temperature zero, blind to the purpose of the study. Two findings are co-equal. First, the score can move with the voice: Gemini 3.1 Pro scores Indian-accented answers 2.5 points lower than the identical answer in an American accent, on a 0 to 100 scale (95% CI -4.06 to -0.94 , $p = 0.005$, $d_z = -0.44$), the only one of the six accent-judge comparisons to reach significance and the only one to survive Benjamini-Hochberg correction across those six tests; Gemini 3.6 Flash shows no such penalty on the same comparison (-0.04 , $p = 1.0$), so two models from the same company, hearing identical audio, disagree about whether an accent lowers the score. Second, the penalty is localized and removable: grading from the model’s own transcript of the same audio erases the effect ($+1.88$, $p = 0.14$), and grading from the frozen gold transcript also erases it (-0.83 , $p = 0.51$), so transcript-first grading is a tested mitigation, not a guess. Fourth, we test five delivery perturbations, filled pauses, faster or slower speech, an 8kHz telephone codec, and background noise, and find no significant score shift under either judge at this sample size, a result we read as inconclusive given the sample rather than as evidence that delivery never leaks into the score. We pre-registered the direction of the accent effect before collecting data, and we also find a partial reversal: Nigerian-accented answers trend toward higher scores under both judges, the opposite of the predicted direction. We report it as we found it.

1 Introduction

An audio language model that grades a spoken answer is being asked to do two things at once: understand the content, and ignore everything about the voice that carries it. Three deployed systems already depend on that second, unstated assumption holding. Automated oral-language assessment tools grade non-native speakers on scales that claim to measure content, not accent (Ma et al., 2025). AI interview-screening pipelines already process hundreds of thousands of real candidate interviews with an LLM-as-judge step in the loop (Allbert et al., 2025), on top of a pre-LLM fairness literature that already catalogued bias sources in automated video interviews (Booth et al., 2023). Speech leaderboards increasingly use an audio-capable model, rather than a human panel, to rank system outputs (Manakul et al., 2025). None of these deployments has been checked against the one manipulation that would prove or disprove the assumption directly: holding what is said perfectly fixed, and changing only who appears to be saying it.




Figure 1: Study design. One frozen spoken answer, rendered in four accents and two genders by TTS, scored blind by audio-input judges. Any score gap between renderings of the same item is a voice effect by construction; measured outcomes previewed at right.

We run that check. We author a bank of 24 short spoken answers once, freeze the text, and render each answer in four accents and two speaker genders using text-to-speech, so that every clip in a comparison carries the identical words. We then have two audio-input judges, Gemini 3.1 Pro and Gemini 3.6 Flash, score every clip on a 0 to 100 rubric, one clip per call, never told that the study concerns voice. Because the text is identical by construction, any score difference between two renderings of the same item can only be explained by the voice.

This paper makes four contributions.

1. A causal audit that puts an audio-input language model in the grader role and varies the evaluatee’s own accent and gender while holding the answer’s content fixed (Section 3).
2. A decomposition design that separates acoustic-channel bias from transcription-channel bias by grading the same item three ways: end-to-end audio, the judge’s own transcript of that audio, and the frozen gold transcript (Section 3.5).
3. A judge-family comparison showing that two models from the same company can disagree about whether an accent changes the score of identical content, one showing a significant penalty and the other showing none (Section 4).
4. A delivery-perturbation battery, filled pauses, speaking rate, telephone-bandwidth audio, and background noise, tested under both a neutral rubric and an explicit instruction to ignore delivery (Section 4.3).

2 What Was Known and What Is New

The premise that audio-language-model judges carry biases is not new. What is new is which bias, and in which role.

AudioJudge (Manakul et al., 2025) establishes that large audio models can act as judges of speech, and catalogs verbosity and position bias in that role. It never varies the identity of the speaker being judged; its evaluatees are text-to-speech systems, not people.

The Voice Behind the Words (Satish et al., 2026a), and its companion study using the same cohort and an explicitly judge-free protocol (Satish et al., 2026b), are the closest papers by topic. It varies the accent and gender of a question-asker across 2,880 controlled interactions and finds that a SpeechLLM’s own reply is rated less helpful for some accents. But the judge in that paper is explicitly text-only and voice-blind: it reads only the assistant’s written reply, never the audio. Voice never reaches the judge in that design. It cannot show, and

does not claim, that an audio-input judge’s score of an evaluatee’s own answer moves with that evaluatee’s voice, which is the exact question this paper asks.

BiasInEar (Wei et al., 2026) varies the voice of a spoken multiple-choice question and measures whether the model’s own answer choice changes. The model here is the test-taker, not a judge; there is no evaluatee and no score assigned to anyone’s spoken answer.

ParaPairAudioBench (Jeon et al., 2026) puts an audio model in the judge role and tests whether it can correctly detect and rank paralinguistic differences such as gender and age. That is the inverse of our question: it asks whether the judge notices the attribute when told to look for it, not whether the attribute leaks into an unrelated content score when the judge is not told to look for it at all.

Bias Analysis of L2 Speaking Assessment Systems (Labroo et al., 2026) audits a real speaking-assessment grader for encoded speaker attributes, but the audit is representational, probing internal activations on naturally varying speakers, not a controlled same-content voice swap.

What remains uncovered by all of the above, and is the subject of this paper, is the exact intersection: an audio-input model in the grader role, the evaluatee’s own voice identity as a randomized treatment, and a content score as the outcome, with the text held identical across every condition.

Three further clusters of recent work sit adjacent to this gap without closing it. A judge-construction literature builds and validates audio judges without a bias axis at all (Wang et al., 2025; Zhang et al., 2026; Chen et al., 2025; Manakul et al., 2025; Park et al., 2026), including judges of speaking style (Chiang et al., 2025) and disfluency robustness of comprehension rather than of scoring (Liu et al., 2025; Takagi et al., 2026). An advisor-role and assistant-role literature varies a user’s paralinguistics and measures the model’s own generated response, decision, or recommendation, not a grade assigned to the user’s speech (Wu et al., 2025; Lin et al., 2026; Satish et al., 2025b;a; Choi et al., 2025). A deployed-assessment literature grades real L2 speech but audits fairness observationally across naturally different speakers rather than through a controlled voice swap on identical content (Uehara, 2026; Ginjala et al., 2026; Koenecke et al., 2020). None of these fifteen papers manipulates an evaluatee’s accent or gender as a randomized treatment on a content score. Miller et al. (2026) comes closest procedurally, holding a transcript fixed while varying affect and prosody to test whether a judge uses audio evidence at all, but the manipulated factor is delivery, not demographic identity, and the outcome is evidence-use correctness rather than a bias in the assigned score.

Two literatures we deliberately stay out of. Voice-agent jailbreak and refusal work studies how prosody and emotion bypass safety training, a different outcome variable than a content grade (Jain et al., 2025). TTS-instrument validity work, on accent fidelity (Zhong et al., 2025; Huang et al., 2026; Xinyuan et al., 2025), on the reliability of TTS-as-stimulus methodology (Yang et al., 2025), and on MOS-annotator demographics (Ren et al., 2026), informs our manipulation-check design but does not audit a judge. We also note Tam and Chen (2025), which finds large modality-driven shifts in a different deployed setting, clinical recommendation, using synthetic voices without an accent axis, and Wang and Sun (2026), which uses the same real-speech corpus we plan as our causal anchor for a diagnostic-feedback task rather than a content-grading one. The human-listener precedent for our exact design predates audio language models entirely: Kang and Rubin (2009) show that identical recorded speech is rated less comprehensible when paired with a different apparent speaker identity, the same causal logic we apply here to an artificial judge instead of a human one.

Significance. If a judge’s score of a spoken answer depends on the speaker’s accent, every deployed pipeline listed in the introduction is partly scoring the speaker rather than the answer. The practical consequence we test directly is whether the standard proposed fix, grading from a transcript instead of from audio, actually removes the effect. Our decomposition arm answers that question for the one bias we find: yes, for that judge, on that accent.

3 Method

3.1 Answer Bank

We authored 24 interview-style question-and-answer pairs across ten everyday domains (science, history, cooking, health, technology, culture, geography, personal finance, daily-life advice, and education), using a frozen text model. Each answer targets a medium quality band, not a strong or a weak one, because a judge has no room to move a score that is already near a ceiling or a floor.

Our first attempt at this calibration failed, and we report the failure because it changed our design. A prompt asking the model for a merely “incomplete or shallow” answer produced answers that a real audio judge scored at a mean of 94.3 out of 100 (standard deviation 4.5, minimum 85), a ceiling with no headroom for any bias to move the score. We diagnosed this before collecting any of the results reported below, and rewrote the generation prompt to require specific, checkable weaknesses in each answer: answering only part of the question, hedging language, one clear factual gap, and a weak ending. We validated the rewritten prompt on five test answers scored by the real judge before regenerating the full bank; the calibration answers scored 35, 35, 45, 60, and 65. The full 24-item bank regenerated under this corrected prompt scored a mean of 61.3 (standard deviation 17.0, range 20 to 90) under Gemini 3.6 Flash, the distribution used throughout this paper.

3.2 Accent and Gender Conditions (Axis A)

Each of the 24 answers was rendered with gemini-3.1-flash-tts-preview in four accents, American, British, Indian, and Nigerian English, crossed with two voices (one designated male, one female), producing 192 clips. The accent is specified as a natural-language style instruction to the TTS model, together with a fixed voice choice per gender. We validated that the TTS model actually renders the requested accent with an 8-clip manipulation check before generating the study data: a separate blind classifier judge correctly identified the intended accent and gender in 8 of 8 clips, all at high stated confidence. Every clip was loudness-normalized (ffmpeg loudnorm) and resampled to 16kHz.

3.3 Delivery Conditions (Axis B)

Using one fixed voice (the American-accented female condition), we generated five delivery perturbations of each of the 24 answers: filled pauses (“um,” “uh” inserted at clause boundaries, re-synthesized by TTS), a slowed and a sped-up rendering (ffmpeg atempo at 0.85 and 1.25), an 8kHz mu-law telephone-bandwidth rendering, and a background-noise condition mixing synthesized pink noise at a measured 10dB signal-to-noise ratio. We do not have a licensed multi-talker babble corpus, so the noise condition uses a pink-noise proxy rather than authentic babble; we state this as a limitation rather than presenting it as equivalent to a babble-noise benchmark.

Every perturbed clip passed a word-error-rate neutrality check before being judged: we transcribed each clip and compared it to the frozen source text, excluding any clip whose transcript changed by more than 5% word error rate, since a perturbation that changes the recognized words is no longer a pure delivery manipulation, the same principle behind channel-specific robustness suites for text pipelines (Hu et al., 2026). Our first transcription pass used a general-purpose instruction and produced an artificially high failure rate, because the transcribing model, an LLM used as an ad hoc speech recognizer, occasionally filled in unrelated caption-like text (video timestamps, a fragment of an unrelated course syllabus) instead of transcribing difficult audio, rather than mis-hearing the actual words. A stricter prompt that explicitly forbade inventing content flipped 28 of 33 initial failures to passing, confirming that most of the initial failures were an artifact of the transcribing model, not of the audio. The final gate excluded 5 of 120 non-baseline conditions as genuine word-error-rate violations. Both prompt arms, a neutral rubric and an explicit instruction to grade content and ignore delivery, were run on the 139 surviving clips.

3.4 Judges and Protocol

Two judges scored every clip: gemini-3.6-flash and gemini-3.1-pro-preview, both called directly through the API at temperature zero, one clip per call, in randomized order, never informed that the study concerns voice or accent. The rubric asks for a single 0 to 100 score covering correctness, completeness, helpfulness, and clarity, plus a one-sentence reason, returned as a structured JSON object. We did not use a third, open-weight judge or an OpenAI judge in this draft; both remain planned additions, discussed in Section 5.

3.5 Decomposition Arm

For a stratified subset of Axis A, the American and Indian accents crossed with both genders across 12 of the 24 items, we graded the same content three ways per judge: end-to-end audio (already collected in Axis A), the judge’s own transcript of that same audio, and the frozen gold text transcript. If a bias found in the audio mode shrinks under the transcript modes, that localizes the bias in the acoustic channel rather than in the model’s general treatment of the content.

3.6 Analysis

The unit of analysis is the within-item score difference between a treatment condition and a reference condition (the American-accent rendering of the same gender for Axis A, the clean baseline for Axis B), computed separately per judge and per condition. We report the mean shift with a bootstrap 95% confidence interval (10,000 resamples), a two-sided sign-flip permutation test (p , 20,000 permutations), and the paired effect size d_z . Every number in Section 4 is read directly from one scripted analysis pass over the raw judgment logs.

4 Results

4.1 Accent and Gender (Axis A)

Table 1 and Figure 2 report the score shift for each non-American accent relative to the American-accent rendering of the same item and the same gender, so the comparison is never confounded by gender. Gemini 3.1 Pro shows one shift that clears our significance threshold: Indian-accented answers score 2.5 points lower than the identical American-accented answer ($p = 0.005$, $d_z = -0.44$). Because this is one of six accent-judge tests, we apply Benjamini-Hochberg correction across all six: the Indian-accent penalty survives at $q = 0.05$ (its $p = 0.005$ falls under the rank-1 critical value of 0.0083), and no other cell does. The same comparison under Gemini 3.6 Flash is indistinguishable from zero (-0.04 , $p = 1.0$). Neither judge shows a significant British-accent effect. Both judges show a Nigerian-accent shift in the positive direction, the opposite of our pre-registered prediction that global-South accents would score lower; neither reaches significance at this sample size. The gender main effect (female minus male, within the American-accent condition) is not significant for either judge, though Gemini 3.6 Flash trends toward scoring the female voice 3.25 points higher ($p = 0.12$).

4.2 Decomposition: Where the Bias Lives

Table 2 traces the Gemini 3.1 Pro Indian-accent penalty across grading modes. The penalty is present and significant when the judge hears audio (-2.50 , $p = 0.006$), reverses in sign and loses significance when the judge grades its own transcript of the identical audio ($+1.88$, $p = 0.14$), and is not significant when the judge grades the frozen gold transcript (-0.83 , $p = 0.51$). This pattern is consistent with the bias living in the acoustic channel specifically, not in a general dispreference for the content once the accent is described in words. Gemini 3.6 Flash, which showed no significant audio-mode effect to begin with, shows a non-significant shift under its own transcript (-4.25 , $p = 0.17$) and under the gold transcript ($+0.83$, $p = 0.36$); we do not read a channel story into a judge that showed no effect to decompose.

Table 1: Score shift, non-American accent minus American accent, same item and gender ($n = 48$ paired comparisons per cell). Bold marks $p < 0.01$.

Accent	Mean shift	95% CI	p
Gemini 3.1 Pro			
UK	-0.42	[-2.08, 1.25]	0.73
Indian	-2.50	[-4.06, -0.94]	0.005
Nigerian	+1.04	[-0.52, 2.60]	0.26
Gemini 3.6 Flash			
UK	-1.08	[-3.90, 1.92]	0.49
Indian	-0.04	[-2.85, 2.73]	1.00
Nigerian	+0.23	[-2.60, 3.08]	0.89




Figure 2: Score shift by accent and judge, 95% bootstrap confidence intervals, $n = 48$ paired comparisons per interval. Right of the dashed line means the accent scored higher than the American baseline. Only the Gemini 3.1 Pro, Indian-accent interval (highlighted) excludes zero at $p < 0.01$.

4.3 Delivery (Axis B)

None of the five delivery perturbations produced a significant score shift under either judge, under either the neutral rubric or the explicit ignore-delivery instruction (all $p > 0.1$; per-condition n ranges from 21 to 24 after the word-error-rate exclusions). Effect sizes were small throughout ($|d_z| < 0.36$). We read this as inconclusive at the current sample size rather than as evidence that delivery never affects content scores: our own budget-driven power analysis at the planning stage flagged that 24 items would leave wide intervals on the secondary axis, and the observed intervals here span both directions of zero in every cell.

5 Limitations

Sample size. The accent and delivery arms use 24 items under a \$20 compute budget, smaller than the 150-item design of the predecessor study this method is adapted from. The Indian-accent effect we report survives that sample size; the null results on the delivery arm and on the British and Nigerian accents may simply be underpowered.

Two judge families, both from one company. Every scored clip was graded by Gemini 3.1 Pro or Gemini 3.6 Flash. We did not have access to an OpenAI audio judge at the time of this draft, and had not yet completed the integration of an open-weight audio judge on Modal. The disagreement between two Gemini models is itself informative, since it shows the effect is not a fixed property of one vendor’s audio pipeline, but a genuine third-family and fourth-family comparison remains future work.

Text-to-speech only, no real-speech anchor yet. Every accent condition in this draft is synthetic. The L2-ARCTIC corpus (Zhao et al., 2018), real recordings of non-native English speakers reading identical prompts against matched native recordings, is the planned causal




Figure 3: The penalty lives in the acoustic channel. The Gemini 3.1 Pro Indian-accent shift under three grading modes; only end-to-end audio shows a significant penalty.

Table 2: Decomposition of the Indian-accent shift by grading mode, Gemini 3.1 Pro, $n = 24$ paired comparisons per mode.

Mode	Mean shift	95% CI	p
Audio	-2.50	[-4.06, -0.94]	0.006
Own transcript	+1.88	[0.00, 3.96]	0.14
Gold transcript	-0.83	[-2.50, 0.83]	0.51

anchor for this project, verified for license and sentence-overlap suitability, but its use here is pending a manual, human-completed registration step with the corpus’s distributor that had not finished at the time of this draft.

Same-family manipulation check. The 8-clip accent-fidelity check used a Gemini model to classify audio generated by a different Gemini model. The classification clips differ in content from the study items, but a same-vendor circularity concern remains until an independent or human-listener check is run.

The noise condition is a proxy. We synthesized pink noise at a measured signal-to-noise ratio rather than mixing in a licensed multi-talker babble corpus; a real babble condition may behave differently.

6 Conclusion

We asked a narrow question with a clean design: if the words are identical, does an audio-input judge’s score still move with the speaker’s voice? For one judge, Gemini 3.1 Pro, on one accent, Indian English, the answer is yes, by 2.5 points on a 100-point scale, and the effect traces to the acoustic channel rather than to the content once transcribed. For a sibling model from the same company, Gemini 3.6 Flash, the answer on the identical comparison is no. Both facts matter for anyone deploying an audio-LLM judge today: which model you pick is not a neutral engineering choice, and grading from a transcript is a concrete, tested mitigation for at least the one failure mode we found. The delivery axis, the real-speech anchor, a third judge family, and a prompt-level mitigation battery are the natural next steps, and are already scoped in this project’s experiment plan.

References

- Rumi Allbert, Nima Yazdani, Ali Ansari, Aruj Mahajan, Amirhossein Afsharrad, and Seyed Shahabeddin Mousavi. Evaluating speech-to-text x llm x text-to-speech combinations for ai interview systems. arXiv preprint arXiv:2507.16835, 2025.
- Brandon M Booth, Louis Hickman, Shree Krishna Subburaj, Louis Tay, Sang Eun Woo, and Sidney K. DMello. Integrating psychometrics and computing perspectives on bias




Figure 4: Delivery perturbations: all null at this sample size. Five perturbations, two judges, two prompt arms; every 95% interval spans zero ($n=21-24$ per cell). Read as inconclusive, not as proof of no effect.

and fairness in affective computing: A case study of automated video interviews. arXiv preprint arXiv:2305.02629, 2023.

Chen Chen, Yuchen Hu, Siyin Wang, Helin Wang, Zhehuai Chen, Chao Zhang, et al. Audio large language models can be descriptive speech quality evaluators. arXiv preprint arXiv:2501.17202, 2025.

Cheng-Han Chiang, Xiaofei Wang, Chung-Ching Lin, Kevin Lin, Linjie Li, Radu Kopetz, et al. Audio-aware large language models as judges for speaking styles. arXiv preprint arXiv:2506.05984, 2025.

Junhyuk Choi, Jihwan Seol, Nayeon Kim, Chanhee Cho, EunBin Cho, and Bugeun Kim. Acoustic-based gender differentiation in speech-aware language models. arXiv preprint arXiv:2509.21125, 2025.

Srishti Ginjala, Eric Fosler-Lussier, Christopher W. Myers, and Srinivasan Parthasarathy. Do llm decoders listen fairly? benchmarking how language model priors shape bias in speech recognition. arXiv preprint arXiv:2604.21276, 2026.

Zizhao Hu, Nathan Elijah Segura, Mohammad Rostami, and Jesse Thomason. Should we type or talk to llm agents? a comprehensive study of voice and keyboard input perturbations. arXiv preprint arXiv:2608.03970, 2026.

Wen-Chin Huang, Nicholas Sanders, and Erica Cooper. Codecmos-accent: A mos benchmark of resynthesized and tts speech from neural codecs across english accents. arXiv preprint arXiv:2603.14328, 2026.

Dhruv Jain, Harshit Shukla, Gautam Rajeev, Ashish Kulkarni, Chandra Khatri, and Shubham Agarwal. Voiceagentbench: Are voice assistants ready for agentic tasks? arXiv preprint arXiv:2510.07978, 2025.

Jisu Jeon, Seungyeon Jwa, Joosung Lee, Jinhyeon Kim, Woojin Chung, Hwiyeol Jo, et al. Parapairaudiobench: Paralinguistic pairwise audio benchmark for lalm-as-a-judge. arXiv preprint arXiv:2606.24648, 2026.

Okim Kang and Donald L Rubin. Reverse linguistic stereotyping: Measuring the effect of listener expectations on speech evaluation. *Journal of Language and Social Psychology*, 28(4):441–456, 2009.

Allison Koenecke, Andrew Nam, Emily Lake, Joe Nudell, Minnie Quartey, Zion Mengesha, Connor Touns, John R Rickford, Dan Jurafsky, and Sharad Goel. Racial disparities in

- automated speech recognition. *Proceedings of the National Academy of Sciences*, 117(14): 7684–7689, 2020.
- Arya Labroo, Mengjie Qian, and Kate Knill. Bias analysis of l2 speaking assessment systems using concept activation vectors. *arXiv preprint arXiv:2608.06300*, 2026.
- Yi-Cheng Lin, Yusuke Hirota, Sung-Feng Huang, and Hung yi Lee. Vibe: Voice-induced open-ended bias evaluation for large audio-language models via real-world speech. *arXiv preprint arXiv:2604.17248*, 2026.
- Hongcheng Liu, Yixuan Hou, Heyang Liu, Yuhao Wang, Yanfeng Wang, and Yu Wang. Vocalbench-df: A benchmark for evaluating speech llm robustness to disfluency. *arXiv preprint arXiv:2510.15406*, 2025.
- Rao Ma, Mengjie Qian, Siyuan Tang, Stefano Bannò, Kate M. Knill, and Mark J. F. Gales. Assessment of l2 oral proficiency using speech large language models. *arXiv preprint arXiv:2505.21148*, 2025.
- Potsawee Manakul, Woody Haosheng Gan, Michael J. Ryan, Ali Sartaz Khan, Warit Siri-chotedumrong, Kunat Pipatanakul, et al. Audiojudge: Understanding what works in large audio model based speech evaluation. *arXiv preprint arXiv:2507.12705*, 2025.
- Kevin Miller, Arjun Chandra, and Venkatesh Saligrama. Do audio language models use paralinguistic evidence? counterfactual audits for response evaluation. *arXiv preprint arXiv:2608.06718*, 2026.
- Joonyong Park, David M. Chan, Yuki Saito, and Hiroshi Saruwatari. Auditing protocol-level shortcuts in large audio language model judges for speech evaluation. *arXiv preprint arXiv:2607.13477*, 2026.
- Wenze Ren, Yi-Cheng Lin, Wen-Chin Huang, Erica Cooper, Ryandhimas E. Zezario, Hsin-Min Wang, et al. Mos-bias: From hidden gender bias to gender-aware speech quality assessment. *arXiv preprint arXiv:2603.10723*, 2026.
- Shree Harsha Bokkahalli Satish, Gustav Eje Henter, and Éva Székely. Do bias benchmarks generalise? evidence from voice-based evaluation of gender bias in speechllms. *arXiv preprint arXiv:2510.01254*, 2025a.
- Shree Harsha Bokkahalli Satish, Harm Lameris, Olivier Perrotin, Gustav Eje Henter, and Éva Székely. Speak your mind: The speech continuation task as a probe of voice-based model bias. *arXiv preprint arXiv:2509.22061*, 2025b.
- Shree Harsha Bokkahalli Satish, Christoph Minixhofer, Maria Teleki, James Caverlee, Ondřej Klejch, Peter Bell, et al. The voice behind the words: Quantifying intersectional bias in speechllms. *arXiv preprint arXiv:2603.16941*, 2026a.
- Shree Harsha Bokkahalli Satish, Maria Teleki, Christoph Minixhofer, Ondřej Klejch, Peter Bell, and Éva Székely. From seeing it to experiencing it: Interactive evaluation of intersectional voice bias in human-ai speech interaction. *arXiv preprint arXiv:2604.13067*, 2026b.
- Masato Takagi, Masaya Kawamura, Reo Shimizu, and Yuma Shirahata. Investigating human-model discrepancies in speech quality assessment via acoustic and prosodic perturbations. *arXiv preprint arXiv:2606.19951*, 2026.
- Zhi Rui Tam and Yun-Nung Chen. Medvoicebias: A controlled study of audio llm behavior in clinical decision-making. *arXiv preprint arXiv:2511.06592*, 2025.
- Eichi Uehara. Interpretable, fairly evaluated automated l2 speaking assessment that beats the single-human ceiling and why pause encoding does not change llm fluency scores. *arXiv preprint arXiv:2608.26137*, 2026.

- Hui Wang, Jinghua Zhao, Yifan Yang, Shujie Liu, Junyang Chen, Yanzhe Zhang, et al. Speechllm-as-judges: Towards general and interpretable speech quality evaluation. arXiv preprint arXiv:2510.14664, 2025.
- Rong Wang and Kun Sun. Prior over evidence: Stereotype-driven diagnosis in llm-based l2 pronunciation feedback. arXiv preprint arXiv:2606.15325, 2026.
- Sheng-Lun Wei, Yu-Ling Liao, Yen-Hua Chang, Hen-Hsen Huang, and Hsin-Hsi Chen. Bias in the ear of the listener: Assessing sensitivity in audio language models across linguistic, demographic, and positional variations. arXiv preprint arXiv:2602.01030, 2026.
- Yihao Wu, Tianrui Wang, Yizhou Peng, Yi-Wen Chao, Xuyi Zhuang, Xinsheng Wang, et al. Evaluating bias in spoken dialogue llms for real-world decisions and recommendations. arXiv preprint arXiv:2510.02352, 2025.
- Henry Li Xinyuan, Zexin Cai, Ashi Garg, Kevin Duh, Leibny Paola García-Perera, Sanjeev Khudanpur, et al. Scalable controllable accented tts. arXiv preprint arXiv:2508.07426, 2025.
- Yifan Yang, Hui Wang, Bing Han, Shujie Liu, Jinyu Li, Yong Qin, et al. Position: Towards responsible evaluation for text-to-speech. arXiv preprint arXiv:2510.06927, 2025.
- Leying Zhang, Bowen Shi, Haibin Wu, Bach Viet Do, and Yanmin Qian. Jastin: Aligning llms for zero-shot audio and speech evaluation via natural language instructions. arXiv preprint arXiv:2605.04505, 2026.
- Guanlong Zhao, Sinem Sonsaat, Alif Silpachai, Ivana Lucic, Evgeny Chukharev-Hudilainen, John Levis, and Ricardo Gutierrez-Osuna. L2-arctic: A non-native english speech corpus. In Interspeech, pages 2783–2787, 2018.
- Jinzuomu Zhong, Suyuan Liu, Dan Wells, and Korin Richmond. Pairwise evaluation of accent similarity in speech synthesis. arXiv preprint arXiv:2505.14410, 2025.